



Chapter 2:

Learn basic Kernel Management

Learning objectives

Upon completing this chapter, the learner should be able to:

- Understand the concept of GRUB
- Boot systems into different runlevels manually
- Interrupt the boot process in order to gain access to a system
- Configure systems to boot into a specific target automatically
- Update the kernel package appropriately to ensure a bootable system
- Manage the kernel modules

Key terms

GRUB2

GRUB menu

grub2-mkconfig

grub.cfg

Linux kernel

Kernel parameters

Emergency target

Rescue target

Kernel version

Kernel image

lsmod modinfo modprobe

Kernel modules

dmesg

journalctl

uname

/proc

Table of content

1. Working with GRUB2	4
1.1. Understanding GRUB2	4
1.2. The /etc/default/grub file	5
1.3. The GRUB menu	8
1.4. Modifying default GRUB2 options	9
1.5. Booting into specific targets	10
2. Linux Kernel	14
2.1. Installed kernel	14
2.2. Kernel directory structure	16
2.2.1. The /boot File System	16
2.2.2. The /proc File System	18
2.2.3. The /lib/modules Directory	20
3. Managing the kernel	22
3.1. Installing and updating	23
3.2. Working with kernel modules	24
3.3. The use of kernel threads and drivers	29
3.4. Analyzing what the kernel is doing	29
3.4.1. The dmesg utility	29
3.4.2. The /proc file system	30
3.4.3. The uname utility	30

1. Working with GRUB2

In RHEL7, a new, enhanced version of GRUB, called [GRUB2](#), has been introduced, replacing the legacy GRUB, it supports both BIOS/MBR and UEFI/GPT combinations. In general, the GRUB2 boot loader is one of the first things that needs to be working well to boot a Linux server, this because after it has been loaded into memory and it takes control, it searches for the kernel, extracts its code into memory, decompresses it, and loads it based on the configuration defined in the [/boot/grub2/grub.cfg](#) file.

1.1. Understanding GRUB2

[GRUB2](#) is installed in the boot sector of your server's hard drive and is configured to load a Linux kernel and the [initramfs](#):

- The [kernel](#) is the heart of the operating system, allowing users to interact with the hardware that is installed in the server.
- The [initramfs](#) contains drivers that are needed to start your server. It contains a mini file system that is mounted during boot. In it are [kernel](#) modules that are needed during the rest of the boot process (for example, the [LVM](#) modules and [SCSI](#) modules for accessing disks that are not supported by default).

Normally, [GRUB2](#) works just fine and does not need much maintenance. In some cases, though, you might have to change its configuration which mainly located in different files in the system, in the following we will learn about these files and how to deal with them.

1.2. The `/etc/default/grub` file

Values defined in this file are used to regenerate the `/boot/grub2/grub.cfg` file, which controls the behavior of GRUB at boot time. Any changes made to the grub file will only take effect after the `grub2-mkconfig` utility has been executed. Here are the default settings from the `/etc/default/grub` file, with an explanation in the table after.

Directive	Description
GRUB_TIMEOUT	Sets the wait time, in seconds, before booting off the default kernel. Default value is 5
GRUB_DISTRIBUTOR	Defines the name of the Linux distribution. “\$(sed 's, release. *\$, g' /etc/system-release)”
GRUB_DEFAULT	Boots the selected option from the previous system boot. saved
GRUB_DISABLE_SUBMENU	Enables/disables the appearance of GRUB submenu. true
GRUB_TERMINAL_OUTPUT	Sets the default terminal. “console”
GRUB_CMDLINE_LINUX	Specifies the command line options to pass to the kernel at boot time. “rd.lvm.lv=vg00/swap vconsole.font=latarcyrheb-sun16 crashker nel=auto vconsole.keymap=us rd.lvm.lv=vg00/root rhgb quiet”
GRUB_DISABLE_RECOVERY	Disables showing system recovery entries in the GRUB menu. true

Generally, you do not need to make any changes to this file, as the default settings are good enough for normal system operation. The [/boot/grub2/grub.cfg](#) file.

This is the main configuration file that controls the behavior of GRUB at boot time. This file is located in the [/boot/grub2](#) directory on BIOS-based systems and in the [/boot/efi/EFI/redhat](#) directory on UEFI-based systems. This file can be regenerated manually with the [grub2-mkconfig](#) utility, or it is automatically regenerated when a new kernel is installed. During this process, any manual changes made to this file are lost.

Here is how you would run this utility to reproduce the [grub.cfg](#) file on BIOS and UEFI systems, respectively:

```
#grub2-mkconfig -o /boot/grub2/grub.cfg  
# grub2-mkconfig -o /boot/efi/EFI/redhat/grub.cfg
```

When this utility runs, it uses the settings defined in the [/etc/default/grub](#) file and in the helper scripts located in the [/etc/grub.d/](#) directory to regenerate this file for kernels located in the [/boot](#) directory.

If a new kernel is added to the system, existing kernel entries will remain in this file and can be chosen in the GRUB menu at startup to boot.

Exercise: Applying Modifications to GRUB2

1. Open the file `/etc/default/grub` with an editor and remove the `rhgb` and `quiet` options from the `GRUB_CMDLINE_LINUX` line.
2. From the same file, set the `GRUB_TIMEOUT` parameter to 10 seconds. Save changes to the file and close the editor.
3. From the command line, type `grub2-mkconfig > /boot/grub2/grub.cfg` to write the changes to GRUB2. (Note that instead of using the redirector `>` to write changes to the `grub.cfg` file, you could use the `-o` option. Both methods have the same result.)
4. Reboot and verify that while booting you see boot messages scrolling by.

1.3. The GRUB menu

The GRUB menu shows a list of bootable kernels at the top, where you can change the selection using the up or down arrow key, or also editing a selected kernel menu entry by:

- Pressing **c** or go to the `grub> command` prompt by pressing **c**. In the edit mode, GRUB loads the selected entry from the `/boot/grub2/grub.cfg` file in an editor, which you are allowed to modify before booting.
- You can press **Ctrl+x** to boot after making the change.
- **Ctrl+c** to switch into the `grub> command` prompt, Or press **ESC** to discard the changes made and go back to the main menu displays one of the entries and the action keys.
- The `grub>` prompt appears when you type **c**. While at the prompt, you can press the **TAB** key to view a list of all available commands that you can run to perform a desired action.

```
Setparams 'Centos Linuxs (3.10.0-1062. 4. 3. e17. x86_64) 7 (core)
    load_video
    set gfxpay load=keep
    insmod gzio
    insmod part_msdos
    insmod xfs
    set root='hd0, msdos1'
    if [x$ feature_platform_search_hint = xy]; then
        Search --no -floppy --fs-uuid -set=root -hint-
bios=hd0, msdos1 -hin\
t-efi= hd0, msdos1 --hint-baremetal=ahci0, msdos1 --hint='hd0,
msdos1' a4818155-b\ a8b-4529-b9ae 2263fa202034
    else
        search --no -floppy --fs -uuid --set=root a4818155-
ba8b-4529-b9ae-2263\ fa202034
    press Ctrl-x to start, Ctrl-c for a command prompt or Escape
to
    discard edits and return to the menu. Pressing Tab lists
possible completions
```

1.4. Modifying default GRUB2 options

If you enter the [GRUB2](#) boot prompt to add kernel startup parameters, the contents of the `/boot/grub2/grub.cfg` file display. From here, you add one-time-only startup options. The following shows the relevant part of the `grub.cfg` file that takes care of loading the Linux kernel. Notice the line that starts with `linux16`, this line specifies all [kernel](#) boot parameters.

Partial Contents of the `/boot/grub2/grub.cfg` Configuration File

```
[root@LP2 ~]# cat /boot/grub2/grub.cfg
menuentry 'CentOS Linux (3.10.0-1062.4.3.el7.x86_64) 7 (Core)'
--class centos --class gnu-linux --class gnu --class os --
unrestricted      $menuentry_id_option      'gnulinux-3.10.0-
862.el7.x86_64-advanced-bf0d0b83-aeaf-4b5f-bb40-2ece311979c0'
{
    load_video
    set gfxpayload=keep
    insmod gzio
    insmod part_msdos
    insmod xfs
    set root='hd0,msdos1'
    if [ x${feature_platform_search_hint} = xy ]; then
        search --no-floppy --fs-uuid --set=root --hint-
bios=hd0,msdos1      --hint-efi=hd0,msdos1      --hint-
baremetal=ahci0,msdos1  --hint='hd0,msdos1'      a4818155-ba8b-
4529-b9ae-2263fa202034
    else
        search --no-floppy --fs-uuid --set=root a4818155-
ba8b-4529-b9ae-2263fa202034
    fi
    linux16      /vmlinuz-3.10.0-1062.4.3.el7.x86_64
root=/dev/mapper/centos-root      ro      crashkernel=auto
rd.lvm.lv=centos/root      rd.lvm.lv=centos/swap      rhgb      quiet
LANG=en_US.UTF-8
    initrd16 /initramfs-3.10.0-1062.4.3.el7.x86_64.img
}
```

To apply modifications to the [GRUB2](#) boot loader, the file [/etc/default/grub](#) is your entry point, do not change the contents of the [/boot/grub2/grub.cfg](#) configuration file directly.

- The most important line in this file is [GRUB_CMDLINE_LINUX](#), which defines how the Linux kernel should be started. In this line, you can apply permanent fixes to the [GRUB2](#) configuration.
- Some likely candidates for removal are the options [rhgb](#) and [quiet](#). These options tell the kernel to hide all output while booting. That is nice to hide confusing messages for end users, but if you are a server administrator, you probably just want to remove these options. Without these you will not have to guess why your server takes a long time after a restart, you'll just be able to see.
- Another interesting parameter is [GRUB_TIMEOUT](#). This defines the amount of time your server waits for you to access the GRUB2 boot menu before it continues booting automatically. If your server runs on physical hardware that takes a long time to get through the BIOS checks, it may be interesting to increase this time a bit.

1.5. Booting into specific targets

RHEL is booted into graphical target state by default. It can also be booted into other non-default, but less capable, operating targets from the GRUB menu. Additionally, in situations when it becomes mandatory to boot the system into an administrative state for carrying out a function that cannot be otherwise performed in other target states or for system recovery, RHEL offers emergency and rescue targets. These special targets can be entered by interacting with the GRUB interface, selecting a boot menu entry, pressing [e](#) to enter the edit mode, and supplying the desired target with the [systemd.unit](#) directive.

For instance, to boot into the emergency target, append

`systemd.unit=emergency.target` (or simply 'emergency') to the default linux kernel line entry, as shown below:

```
linux16 /vmlinuz-3.10.0-123.el7.x86_64
root=UUID=964201bb-1e32-4794-a2f2-7a33e2fb591a ro
rd.lvm.lv=vag00/swap
vconsole.font=latacyrheb-sun16 crashke\rnel=auto
vconsole.keymap=us rd.lvm.lv=vag00/root rhgb quiet
systemd.unit=emer\gency.target_
```

Press Ctrl+x after making the modification to boot the system into the supplied target. You will be required to enter the root user password to log on. Run `systemctl reboot` after you are done to reboot the system.

Similarly, you can enter `systemd.unit=rescue.target` (or simply 1, s, or single) with the linux kernel line entry and press Ctrl+x to boot into the rescue target, which is also referred to as the single-user mode.

For this exercise, assume that the root user password has been lost or forgotten, and it needs to be reset. You will boot the system into a special shell in order to reset the root password.

Exercise: Resetting the root user password

1. Reboot or reset server1, and interact with GRUB by pressing a key before the autoboot timer runs out. Highlight the default kernel entry in the GRUB menu and press e to enter the edit mode. Scroll down and you will find a boot string similar to the following:

```
linux16 /vmlinuz-3.10.0-1062.4.3.el7.x86_64
root=/dev/mapper/centos-root ro crashkernel=auto
rd.lvm.lv=centos/root rd.lvm.lv=centos/swap rhgb quiet\
LANG=en_US.UTF-8
initrd16 /initramfs-3.10.0-1062.4.3.el7.x86-64.img
```

2. Modify this kernel string and append “init = /sysroot/bin/sh” to the end of the line to look like:

```
linux16 /vmlinuz-3.10.0-1062.4.3.el7.x86_64
root=/dev/mapper/centos-root ro crashkernel=auto
rd.lvm.lv=centos/root rd.lvm.lv=centos/swap rhgb quiet\
LANG=en_US.UTF-8\
init=/sysroot/bin/sh
initrd16 /initramfs-3.10.0-1062.4.3.el7.x86-64.img
```

3. Press Ctrl+x when done to boot to the special shell. The system mounts the root file system read-only on the /sysroot directory. Make /sysroot appear as mounted on / using the chroot command:

```
# chroot /sysroot
```

4. Remount the root file system in read/write mode with the mount command:

```
# mount -o remount,rw /
```

5. Enter a new password for root by invoking the passwd command:

```
# passwd
```

6. Create an empty, hidden file called `.autorelabel` at the root of the directory tree to instruct the system to perform SELinux relabeling upon next reboot:

```
# touch /. autorelabel
```

7. Exit out of the special shell:

```
# exit
```

8. Reboot the system:

```
# reboot
```

2. Linux Kernel

The Linux [kernel](#) is the heart of the operating system, it is the layer between the user who works with Linux from a shell environment and the hardware that is available in the computer on which the user is working. It is doing so by managing the I/O instructions it receives from the software and translating those to processing instructions that are to be executed by the central processing unit and other hardware in the computer. The [kernel](#) also takes care of handling essential operating system tasks like the scheduler, which makes sure that processes that are started on the operating system are handled by the CPU.

2.1. Installed kernel

The default [kernel](#) installed during the installation is usually adequate for most system needs; however, it requires a rebuild when a new functionality is added or removed. The new functionality may be introduced by installing a new [kernel](#), upgrading the existing one, installing a new hardware device, or changing a critical system component.

RHEL allows us to generate and store several custom kernels with varied configuration and required modules, but only one of them is active at a time and the others may be loaded via GRUB. To list all installed kernel packages:

```
[root@LP2 ~]# yum installed kernel-*  
Installed Packages  
kernel.x86_64          3.10.0-862.el7        @anaconda  
kernel.x86_64          3.10.0-1062.4.3.el7    @updates  
kernel-tools.x86_64    3.10.0-862.el7        @anaconda  
kernel-tools-libs.x86_64 3.10.0-862.el7 @anaconda
```

Likewise, an existing functionality that is no longer required may be removed to make the kernel smaller, resulting in improved performance and reduced memory utilization.

To control the behavior of the modules, and the [kernel](#) in general, several tunable parameters are set. It defines a baseline for kernel functionality. Some of these parameters must be tuned to allow certain applications and database software to be installed smoothly and operate properly.

Determining Kernel Version

To determine the version of the running [kernel](#) on the system, run the `uname` command:

```
[root@LP2 ~]# uname -r  
3.10.0-1062.4.3.el7.x86_64
```

The output indicates the kernel version currently in use is 3.10.0-1062.4.3.el7.x86_64

2.2. Kernel directory structure

Kernel and its support files are stored at different locations in the directory hierarchy, of which three locations: `/boot`, `/proc`, and `/lib/modules` —are of significance and are explained below.

2.2.1. The `/boot` File System

The `/boot` file system is created at system installation and its purpose is to store kernel and associated files. This file system also stores any updated or modified kernel data.

Listing the content of the `/boot` produces the following information:

```
[root@lp2 ~]# ls /boot/
config-3.10.0-1127.19.1.el7.x86_64
config-3.10.0-1127.el7.x86_64
efi
grub
grub2
initramfs-0-rescue-cc7023446ea74f37ab8014f3e85845c5.img
initramfs-3.10.0-1127.19.1.el7.x86_64.img
initramfs-3.10.0-1127.19.1.el7.x86_64kdump.img
initramfs-3.10.0-1127.el7.x86_64.img
initramfs-3.10.0-1127.el7.x86_64kdump.img
symvers-3.10.0-1127.19.1.el7.x86_64.gz
symvers-3.10.0-1127.el7.x86_64.gz
System.map-3.10.0-1127.19.1.el7.x86_64
System.map-3.10.0-1127.el7.x86_64
vmlinuz-0-rescue-cc7023446ea74f37ab8014f3e85845c5
vmlinuz-3.10.0-1127.19.1.el7.x86_64
vmlinuz-3.10.0-1127.el7.x86_64
```

The output indicates that:

- the current kernel is [vmlinuz-3.10.0-1062.4.3.el7.x86_64](#)
- its boot image is stored in the [initramfs-3.10.0-1062.4.3.el7.x86_64.img](#)
- and configuration in the [config-3.10.0-1062.4.3.el7.x86_64](#)
- A sub-directory [/boot/grub2](#) contains GRUB information as shown below:

```
[root@lp2 ~]# ll /boot/grub2/
total 32
-rw-r--r--. 1 root root   84 Oct 27 19:22 device.map
drwxr-xr-x. 2 root root   25 Oct 27 19:22 fonts
-rw-r--r--. 1 root root 5218 Nov 17 23:50 grub.cfg
-rw-r--r--. 1 root root 1024 Nov 17 23:50 grubenv
drwxr-xr-x. 2 root root 8192 Oct 27 19:22 i386-pc
drwxr-xr-x. 2 root root 4096 Oct 27 19:22 locale
```

The key file in [/boot/grub2](#) is [grub.cfg](#), which maintains a list of available kernels and defines the default kernel to boot, along with other information.

2.2.2. The /proc File System

`/proc` is a virtual file system and its contents are created in memory at system boot and destroyed when the system goes down. Underneath this file system lie current hardware configuration and status information.

Listing the content of the `/proc` produces several files and sub-directories:

- Some sub-directory names are numerical and contain information about a specific process, with process ID matching the sub-directory name. Within each sub-directory, there are files and further sub-directories that include information, such as memory segment specific to that particular process.

```
[root@lp2 ~]# ls /proc/
1      14      1812    1927    1982    2615    292     378     501     615     701     731
asound      filesystems  locks      self      vmstat
10     1484    1816    1928    1986    2620    294     379     51      616     702     732
buddyinfo   fs          mdstat     slabinfo   zoneinfo
103     1489    1840    1932    1989    27       295     38      510     618     703     733
bus         interrupts  meminfo    softirqs
11      1490    1848    1936    2       2732    296     391     52      624     704     734
cgroups     iomem      misc       stat
1118     1586    1865    1938    20      2779    299     392     53      625     708     735
cmdline     ioports    modules    swaps
1120     16     1886    1942    2004    279     300     4       6       629     709     741
consoles    irq        mounts     sys
1122     1729    1889    1944    2031    2796    3114    46      607     630     710     744
cpuinfo     kallsyms   mpt        sysrq-trigger
1130     1730    1892    1948    21      28      3262    476     608     66      712     790
crypto      kcore      mtrr       sysvipc
1136     1790    1894    1953    22      280     3316    477     609     673     714     8
devices     keys       net        timer_list
1139     1794    19      1959    23      281     3386    478     610     674     716     800
diskstats   key-users  pagetypeinfo timer_stats
1140     18     1906    1962    24      282     3389    48      611     676     718     832
dma         kmsg      partitions  tty
1141     1800    1916    1965    25      286     35      49      612     678     720     865
driver      kpagecount sched_debug uptime
12      1801    1920    1972    26      287     36      5       613     680     725     9
execdomains kpageflags schedstat   version
13      1807    1926    1978    2603    29      37      50      614     7       729     acpi
fb         loadavg    scsi       vmallocinfo
```

- Other files and sub-directories contain configuration data for system components. If you wish to view configuration for a particular item, such as the CPU or memory, cat the contents of [cpuinfo](#) and [meminfo](#) files as shown below:

```
[root@lp2 ~]# cat /proc/cpuinfo
processor       : 0
vendor_id     : GenuineIntel
cpu family    : 6
model         : 142
model name    : Intel(R) Core(TM) i5-7200U CPU @ 2.50GHz
stepping      : 9
microcode     : 0xb4
cpu MHz       : 2712.006
cache size    : 3072 KB
physical id   : 0
...
processor       : 1
vendor_id     : GenuineIntel
cpu family    : 6
model         : 142
model name    : Intel(R) Core(TM) i5-7200U CPU @ 2.50GHz
stepping      : 9
microcode     : 0xb4
cpu MHz       : 2712.006
cache size    : 3072 KB
physical id   : 2
...
...
[root@lp2 ~]# cat /proc/meminfo
MemTotal:      3861308 kB
MemFree:       3023192 kB
MemAvailable:  3177028 kB
Buffers:       30124 kB
Cached:        324548 kB
SwapCached:    0 kB
Active:        340732 kB
Inactive:      281980 kB
...
```

The data stored under [/proc](#) is referenced by a number of system utilities, including [top](#), [ps](#), [uname](#), and [vmstat](#), to display information.

2.2.3. The `/lib/modules` Directory

This directory holds information about kernel modules. Underneath it are located sub-directories specific to the kernels installed on the system. For example, listing the content of `/lib/modules` below shows that there is only one kernel on this system:

```
[root@lp2 ~]# ll -h /lib/modules
total 8.0K
drwxr-xr-x. 7 root root 4.0K Nov 17 23:50 3.10.0-1127.19.1.el7.x86_64
drwxr-xr-x. 7 root root 4.0K Oct 27 19:22 3.10.0-1127.el7.x86_64
[root@lp2 ~]#
```

Now listing the content of the kernel sub-directory will show:

```
[root@lp2 ~]# ll -h /lib/modules/3.10.0-1127.19.1.el7.x86_64/
total 3.3M
lrwxrwxrwx. 1 root root 44 Nov 17 23:48 build ->
/usr/src/kernels/3.10.0-1127.19.1.el7.x86_64
drwxr-xr-x. 2 root root 4.0K Aug 25 20:27 extra
drwxr-xr-x. 12 root root 4.0K Nov 17 23:48 kernel
-rw-r--r--. 1 root root 840K Nov 17 23:50 modules.alias
-rw-r--r--. 1 root root 801K Nov 17 23:50 modules.alias.bin
-rw-r--r--. 1 root root 1.4K Aug 25 20:28 modules.block
-rw-r--r--. 1 root root 7.3K Aug 25 20:27 modules.builtin
-rw-r--r--. 1 root root 9.3K Nov 17 23:50 modules.builtin.bin
-rw-r--r--. 1 root root 267K Nov 17 23:50 modules.dep
-rw-r--r--. 1 root root 374K Nov 17 23:50 modules.dep.bin
-rw-r--r--. 1 root root 361 Nov 17 23:50 modules.devname
-rw-r--r--. 1 root root 140 Aug 25 20:28 modules.drm
-rw-r--r--. 1 root root 69 Aug 25 20:28 modules.modetesting
-rw-r--r--. 1 root root 1.8K Aug 25 20:28 modules.networking
-rw-r--r--. 1 root root 96K Aug 25 20:27 modules.order
-rw-r--r--. 1 root root 569 Nov 17 23:50 modules.softdep
-rw-r--r--. 1 root root 392K Nov 17 23:50 modules.symbols
-rw-r--r--. 1 root root 482K Nov 17 23:50 modules.symbols.bin
lrwxrwxrwx. 1 root root 5 Nov 17 23:48 source -> build
drwxr-xr-x. 2 root root 4.0K Aug 25 20:27 updates
drwxr-xr-x. 2 root root 4.0K Nov 17 23:48 vdso
drwxr-xr-x. 3 root root 4.0K Nov 17 23:49 weak-updates
[root@lp2 ~]#
```

There are several files and a few sub-directories here those hold module-specific information. One of the key sub-directories is

[lib/modules/3.10.0-1062.4.3.el7.x86_64/kernel/drivers](#), which stores modules categorized in groups in various sub-directories as shown in the listing below:

```
[root@lp2 ~]# ls /lib/modules/3.10.0-1127.19.1.el7.x86_64/kernel/drivers/
acpi      bluetooth  dax        gpio       i2c        isdn
mfd      nvdimmm   platform  rtc        tty        video
ata       cdrom     dca        gpu        idle       leds
misc     nvme      power     scsi       uio        virtio
auxdisplay char      dma        hid        iio        md
mmc      parport   powercap   ssb        usb        watchdog
base     cpufreq   edac       hv          infiniband media
mtd      pci       pps        staging     uwb        xen
bcma     cpuidle   firewire   hwmon       input      memstick
net      pcmcia    ptp        target      vfio
block    crypto    firmware   hwtracing   iommu      message
ntb      pinctrl   pwm        thermal     vhost
```

Several module categories exist, such as ata, bluetooth, cdrom, firewire, input, net, pci, scsi, usb, and video. These categories contain modules to control the hardware components associated with them.

3. Managing the kernel

Managing the [kernel](#) involves performing several tasks, such as installing and updating the [kernel](#), and listing, displaying, loading, unloading, and installing modules. It also includes the task of adding and removing modules to the initial ram disk; however, this is not usually done manually. This task is accomplished automatically when a new [kernel](#) is rebuilt.

- The tools that are used to install and update the kernel are the yum and rpm commands
- While those for managing modules are [lsmod](#), [modinfo](#), [modprobe](#), and [depmod](#). The module management tools are part of the [kmod](#) package, and they automatically take into account any dependent modules during their execution

3.1. Installing and updating

Unlike handling other package installs and upgrades, installing and updating kernel packages require extra care as you might end up leaving your system in an unbootable state. The new kernel addresses existing issues as well as adds bug fixes, security updates, new features, and support for additional hardware devices.

When you upgrade the Linux [kernel](#), a new version is installed and used as the default one. The old version of the kernel file will still be available, though. This ensures that your computer can still boot if in the new kernel nonsupported functionality is included. To install a new version of the kernel, you can use the command [yum upgrade kernel](#) or [yum install kernel](#) command where both commands install the new kernel besides the old one.

The kernel files for the last four kernels that you have installed on your server will be kept in the [/boot](#) directory. The [GRUB2](#) boot loader automatically picks up all kernels that it finds in this directory.

Exercise: Install a new Kernel

1. Run the `uname` command and check the version of the running kernel
2. Run the `yum` command to install the latest available kernel from the subscription management service using either the `update` or the `install` subcommand
3. Confirm that the kernel package has been updated
4. The `/boot/grub2/grub.cfg` file gets the newest kernel
5. Reboot the system and you will see the new kernel menu entry shows up in the GRUB boot list. The system will boot with this kernel as the install process has marked it as the default kernel
6. Run the `uname` command again after the reboot to confirm the loading of the new kernel

3.2. Working with kernel modules

In the old days of Linux, kernels had to be compiled to include all drivers that were required to support computer hardware. Other specific functionality needed to be compiled into the kernel as well. Since the release of Linux kernel 2.0 in the mid-1990s, kernels are no longer compiled but modular. A modular kernel consists of a relatively small core kernel and provides driver support through modules that are loaded when required. Modular kernels are very efficient, as only those modules that are really needed are included.

Linux [kernel](#) modules normally are loaded automatically for the devices that need them, but you will sometimes have to load the appropriate [kernel](#) modules manually. A few commands are used for manual management of kernel modules. The following table provides an overview.

An alternative method of loading kernel modules is by doing this through the [/etc/modules-load.d](#) directory. In this directory, you can create files to load modules automatically that are not loaded by the udev method already.

Command	Use
Lsmod	Lists currently loaded kernel modules.
Modinfo	Displays information about kernel modules.
Modprobe	Loads kernel modules, including all of their dependencies.
modprobe -r	Unloads kernel modules, considering kernel module dependencies.

The first command to use when working with kernel modules is `lsmod`, it lists all [kernel](#) modules that are currently used, including the modules by which this specific module is used. The command `lsmod | head` shows the output of the first 10 lines of the `lsmod` command.

```
[root@lp2 ~]# lsmod | head
Module                Size  Used by
xt_CHECKSUM           12549  1
ipt_MASQUERADE        12678  3
nf_nat_masquerade_ipv4 13463  1 ipt_MASQUERADE
tun                   36164  1
devlink               60067  0
ip6t_rpfilter         12595  1
ip6t_REJECT           12625  2
nf_reject_ipv6        13717  1 ip6t_REJECT
ipt_REJECT            12541  4
[root@lp2 ~]#
```

If you want to have more information about a specific kernel module, you can use the `modinfo` command. This gives complete information about the specific kernel modules, including two interesting sections: the alias and the parms. A module alias is another name that can also be used to address the module. The following shows partial output of the `modinfo e1000` command.

```
[root@lp2 ~]# modinfo xt_CHECKSUM
filename:                               /lib/modules/3.10.0-
1127.19.1.el7.x86_64/kernel/net/netfilter/xt_CHECKSUM.ko.xz
alias:                                  ipt_CHECKSUM
alias:                                  ip6t_CHECKSUM
description:                            Xtables: checksum modification
author:                                  Michael S. Tsirkin <mst@redhat.com>
license:                                  GPL
retpoline:                               Y
rhelversion:                             7.8
srcversion:                              188D3664B303A3171F4928C
depends:
intree:                                  Y
vermagic:                                3.10.0-1127.19.1.el7.x86_64 SMP mod_unload
modversions
signer:                                  CentOS Linux kernel signing key
sig_key:
B1:6A:91:CA:C9:D6:51:46:4A:CB:7A:D9:B8:DE:D5:57:CF:1A:CA:27
sig_hashalgo:                            sha256
[root@lp2 ~]#
```

To manually load and unload modules, you can use `modprobe` and `modprobe -r` commands. On earlier Linux versions, you may have used the `insmod` and `rmmod` commands. These should be used no longer because they do not consider kernel module dependencies.

Exercise: Managing kernel modules from the command line

1. Open a root shell and type `lsmod | head`. This shows all kernel modules currently loaded.
2. Type `modprobe ext4` to load the ext4 kernel module. Verify that it is loaded, using the `lsmod` command again.
3. Type `modinfo ext4` to get information about the ext4 kernel module. Notice that it does not have any parameters.
4. Type `modprobe -r ext4` to unload the ext4 kernel module again.
5. Type `modprobe -r xfs` to try to unload the xfs kernel module. Notice that you get an error message as the kernel module is currently in use.

Managing kernel module parameters

You might sometimes want to load [kernel](#) modules with specific parameters. If this is the case, you first need to find out which parameter you want to use. If you have found the parameter you want to use, you can load it manually, specifying the name of the parameter followed by the value that you want to assign. To make this an automated procedure, you can create a file in the [/etc/modprobe.d](#) directory, where the module is loaded including the parameter you want to be loaded.

In the following Exercise, you see how to do this using the `cdrom` kernel module.

Exercise: Loading Kernel Modules with Parameters

1. Type `lsmod | grep cdrom`. If you have used the optical drive in your computer, this module should be loaded, and it should also indicate that it is used by the `sr_mod` module.
2. Type `modprobe -r cdrom`. This will not work because the module is in use by the `sr_mod` module.
3. Type `modprobe -r sr_mod, modprobe -r cdrom`. This will unload both modules.
4. Type `modinfo cdrom`. This will show information about the `cdrom` module including the parameters that it supports. One of these is the `debug` parameter, which supports a Boolean as its value.
5. Now use the command `modprobe cdrom debug=1`. This will load the `cdrom` module with the `debug` parameter set to on.
6. Type `dmesg`. For some kernel module, load information is written to the kernel ring buffer which can be displayed using the `dmesg` command. Unfortunately, this is not the case for the `cdrom` kernel module.
7. Create a file with the name `/etc/modprobe.d/cdrom` and give it the following contents: `options cdrom debug=1`, This will enable the parameter every time the `cdrom` kernel module will be loaded.

3.3. The use of kernel threads and drivers

The operating system tasks that are performed by the kernel are implemented by different kernel threads. Kernel threads are easily recognized with a command like `ps aux`. The kernel thread names are listed between square brackets.

Another important task of the Linux kernel is hardware initialization. To make sure that this hardware can be used, the Linux kernel uses drivers. Every piece of hardware contains specific features, and to use these features, a driver must be loaded. The Linux kernel is modular, and drivers are loaded as kernel modules.

3.4. Analyzing what the kernel is doing

To help analyze what the kernel is doing, some tools are provided by the Linux operating systems `dmesg` utility, `/proc` file system and the `uname` utility.

3.4.1. The `dmesg` utility

The first utility to consider whether detailed information about the `kernel` activity is required is `dmesg`. This utility shows the contents of the kernel ring buffer, an area of memory where the Linux `kernel` keeps its recent log messages.

In the `dmesg` output, all kernel-related messages are shown. Each message starts with a time indicator that shows at which specific second the event was logged. This time indicator is relative to the start of the `kernel`, which allows you to see exactly how many seconds have passed between the start of the `kernel` and a particular event. (Notice that the `journalctl -k / --dmesg` commands show clock time, instead of time that is relative to the start of the `kernel`.) This time indicator gives a clear indication of what has been happening and at which time it has happened.

3.4.2. The `/proc` file system

Another valuable source of information is the `/proc` file system. The `/proc` file system is an interface to the Linux `kernel`, and it contains files with detailed actual status information on what is happening on your server. Many of the performance related tools mine the `/proc` file system for more information.

As an administrator, you will find that some of the files in `/proc` are very readable and contain actual status information about CPU, memory, mounts, and more.

Take a look, for instance, at `/proc/meminfo`, which gives detailed information about each memory segment and what exactly is happening in these memory segments.

3.4.3. The `uname` utility

A last useful source of information that should be mentioned here is the `uname` command. This command gives different kinds of information about your operating system. Type, for instance, `uname -a` for an overview of all relevant parameters of `uname -r` to see which kernel version currently is used. This information also shows when using the `hostnamectl status` command.

Quiz

Chapter review questions

1. The systemd command may be used to rebuild a new kernel.
 - a. True
 - b. False
2. What is the location of the grub. efi file in the UEFI-based systems?
 - a. /boot/efi/EFI/redhat
 - b. /boot/grub2/grub. cfg
 - c. /bin/efi/EFI/redhat
3. What is the location of the grub. cfg file in the BIOS-based systems?
 - a. /boot/efi/EFI/redhat
 - b. /boot/grub2/grub.cfg
 - c. /bin/efi/EFI/redhat
4. Which file stores the location information of the boot partition on the BIOS systems?
 - a. /boot/bios
 - b. grub. conf
 - c. /proc/partition
5. Which file would we want to view to obtain processor information?
 - a. /proc/cpuinfo
 - b. /proc/meminfo
 - c. /proc/partition

6. Which file would we want to view to obtain memory information?
- a. /proc/cpuinfo
 - b. /proc/meminfo
 - c. /proc/partition
7. Which command can we use to determine the kernel release information?
- a. Kernel-release
 - b. dmesg
 - c. uname -r
8. The lsmod command is used to rebuild modules.
- a. True
 - b. False
9. Which command can we use to unload a module?
- a. modprobe
 - b. lsmod
 - c. modinfo
10. We cannot use the yum command to upgrade a Linux kernel. True or False?
- a. True
 - b. False

- 11.** Which configuration file should you modify to apply common changes to GRUB2?
- a. `/boot/grub2/grub.cfg`
 - b. `/etc/grub.cfg`
 - c. `/grub2/grub.config`
- 12.** After applying changes to GRUB 2 configuration, which command should you run?
- a. `modprobe`
 - b. `mkconfig`
 - c. `grub2-mkconfig`
- 13.** Which command shows a list of kernel modules that are currently loaded?
- a. `lsmod`
 - b. `modprobe`
 - c. `modinfo`
- 14.** How do you install a new version of the kernel?
- a. `yum update`
 - b. `yum update kernel`
 - c. `yum kernel install`
- 15.** Which command enables you to discover kernel module parameters?
- a. `lsmod`
 - b. `modinfo`
 - c. `modprob`

16. How do you unload a kernel module?

- a. `modprobe -r`
- b. `lsmod -r`
- c. `modinfo`

17. It contains drivers that are needed to start your server

- a. `initramfs`
- b. `grub2`
- c. `systemd`

18. Setting the `GRUB_TIMEOUT` parameter to 10, will?

- a. Restart the system after 10 seconds.
- b. Boot the system into the default kernel after 10 seconds.
- c. Shutdown the system after 10 seconds.

19. You can apply permanent fixes to the GRUB 2 configuration using the option:

- a. `GRUB_CMDLINE_LINUX`
- b. `GRUB_TIMEOUT`
- c. `GRUB_RHGB`

20. To list all installed kernels use:

- a. `lsmod kernel-*`
- b. `yum display kernels`
- c. `yum installed kernel-*`

Answers to chapter Review Questions

1. b

2. a

3. b

4. b

5. a

6. b

7. c

8. b

9. a

10. b

11. a

12. c

13. a

14. b

15. b

16. a

17. a

18. b

19. a

20. c